

M. TALHA SIDDIQUI

Lead AI Engineer | Agentic Systems · Multi-Agent Orchestration · RAG · Azure AI



About

- ✉ mtalha.dev@protonmail.com
- ☎ (0092) 3122853636
- 💻 mtalhas.github.io
- 🌐 in/mtalhas
- 🗣️ mtalhas
- 💬 iaim.io · live AI assistant
- ☁️ Azure Marketplace · AI offerings

AI & LLM

Advanced

- LangGraph
- Multi-Agent Orchestration
- RAG
- LLM Evaluation
- Azure OpenAI
- MCP

Cloud & DevOps

Advanced

- Azure
- Azure AI Foundry
- Terraform / Bicep
- Docker
- Container Apps
- CI/CD

Development

Expert

- Python
- C# / .NET
- Go
- TypeScript
- FastAPI

Architecture & Security

Advanced

- RLS Multi-Tenancy
- OAuth2 / OIDC
- CIAM (Entra)
- HSM Signing
- LLMOps

Certifications

- Certified Expert, API Management
Broadcom, 2023
- Proven Professional, API Management
Broadcom, 2023
- Proven Professional, API Management
Broadcom, 2018

Summary

Hands-on Lead AI Engineer with 10+ years in software engineering and cloud architecture, specializing in production language-model systems. Designs and ships multi-agent orchestration, RAG pipelines, and regulated-cloud AI platforms end to end, and built the AI practice that productized them into public Azure Marketplace offerings. Operates a live client-facing AI assistant in production. Brings engineering depth to build and the communication skill to drive adoption across stakeholders and clients.

Experience

Architecture In Motion Inc.

09/2025 - Present

Lead AI Engineer (Manager, AI/LLM)

Hands-on technical lead and AI practice owner across a 6-15 person team; ships production code while setting architecture and strategy.

- Primary engineering contributor across the company's AI platform org, owning architecture and shipping production code in the agent backend and the customer identity (CIAM) layer
- Operate and own the company's production AI assistant, live at iaim.io: a custom Azure-native build (Azure Functions, Azure AI Search for RAG, Azure OpenAI) with an XState v5 conversation engine and a 15-alert monitoring stack with synthetic tests and an IT runbook
- Architected a regulated healthcare AI compliance platform using LangGraph multi-agent orchestration on Azure, with tamper-evident cryptographic evidence and standards-based interoperable output
- Built and led a multi-agent assistant platform across dev, QA, and demo environments with a connector and OAuth framework and a comprehensive automated test suite
- Productized the AI practice into two published Azure Marketplace offerings (AI Assessment and AI Implementation)
- Established Infrastructure as Code (Terraform and Bicep) and Azure DevOps CI/CD pipelines from the ground up
- Led a 6-15 person AI team and Microsoft and IBM partner engagements; authored executive briefings and go-to-market materials; ran client demos and a healthcare showcase

Architecture In Motion Inc.

08/2017 - 08/2025

Senior DevOps Engineer

8-year tenure spanning API management, cloud infrastructure, and the build-out of the company's first production AI platforms.

- Deployed enterprise-grade API Gateway solutions across finance and transportation clients, implementing OAuth2 and OIDC security protocols
- Built CI/CD pipelines in Azure DevOps with infrastructure as code across multiple client environments
- Managed containerized deployments on Docker and Kubernetes with monitoring via Azure Monitor and Application Insights
- Established testing frameworks using Postman, JMeter, and Azure Test Plans for performance and security validation

Architecture In Motion Inc.

08/2016 - 08/2017

API-M Support Consultant

- Engineered enterprise-grade REST API security policies using Layer7 API Gateway
- Automated routine operational tasks through Bash scripting and API Gateway policy development
- Managed Linux environments and MySQL/Oracle databases maintaining 99.9% uptime across Production and Lower environments

Education

MS Computer Science (MSCS)
MAJU, 2017 - 2018

MS Computer Science (MCS)
MAJU, 2015 - 2016

BS Polymer & Petrochemical Engineering
NEDUET, 2010 - 2013

- Implemented advanced log analysis using ELK Stack and Lucene queries in Kibana for rapid root cause analysis

GeekUnit Private Limited
Frontend Web Developer

01/2015 - 08/2015

- Built responsive web interfaces for multiple clients using HTML, CSS, and JavaScript
- Developed reusable component library and collaborated with UI/UX designers and backend developers
- Optimized web page load times through efficient coding practices and image optimization

Projects

AI Portfolio Assistant — mtalhas.github.io (live)

An AI assistant I designed and built entirely on my own, end to end, and the chatbot you can try on this site. It runs in production at almost no cost (it understands language on its own instead of paying per call to AI services), and it is fully secured and deployed. Proof that I can take an AI product from idea to live, cost-aware production by myself.

LLM Evaluation Framework

A framework I built to systematically test and quality-check the answers AI models produce, scoring them against clear rubrics and stress-testing them with adversarial checks across multiple models. Evaluation is the hardest, most overlooked part of shipping reliable AI, and this is how I keep mine trustworthy. Private; walkthrough on request.